

# Type Checking and Type Inference for Quack

Version 0.2, October 2018

Michal Young

## 1 Introduction

After several attempts to write a nice, succinct description of the type system of Quack as one short section of the Quack manual, this document represents my surrender to the reality that it really needs its own document. This is not because the type system is complicated. Quack’s type system is very simple compared to the type systems and type inference in many contemporary programming languages. Nonetheless explaining it clearly enough for implementation by first-time compiler writers, and clarifying it in my own mind, requires a bit more than a short manual section.

## 2 Recap of Relevant Quack Features

With a name like Quack, one might expect duck typing, a variety of dynamic typing. In a *dynamically* typed language, values have types (represented at run-time by tags), but variables don’t. In a *statically* typed language, a type is associated with each variable. Values may also have types, but the dynamic type of an object must conform to the static type of the variable that holds it. A *strongly* static typed language is one in which this conformance relation is verified at translation time.

Quack is a statically typed object-oriented language with single inheritance. There are no ‘primitive’, unboxed<sup>1</sup> types as in Java. Quack’s *Int* is like Java’s *Integer* and not like Java’s *int*. *Obj* is the root of the class hierarchy.

The usual relation holds between the class hierarchy and the type hierarchy in Quack: Every class is a type, and (almost) every type is a class. *D* extends *C* in the class hierarchy iff  $D <: C$  in the type system. The subtype relation is a partial order, which means it is reflexive, antisymmetric, and transitive. For any types  $A, B, C$

$$\frac{}{A <: A} \quad , \quad \frac{A <: B, B <: C}{A <: C} \quad , \quad \frac{A <: B, B <: A}{A = B}$$

---

<sup>1</sup>We call a type like *Integer* in Java a “boxed” type, because the *Integer* object is like a box around the *int* value it holds. When we want to use an *int* or a *char* or a *boolean* value in a collection like *List* or *HashMap*, which expects some kind of *Object*, we must “box” it. The primitive types *int*, *boolean*, and *char* are called “primitive” or “unboxed” types because they are not themselves objects. A language with unboxed types may gain a performance advantage, both because a primitive type requires less storage and because it can often be allocated in the program stack, while objects are allocated in the program heap (always in Java, only sometimes in C++). This performance advantage comes at an expense in complexity for the programmer and for the compiler writer.

The correspondence between types and classes is almost complete, but for purposes of the formal system we will introduce two special types that do not correspond to classes. The special type  $\top$  lies above *Obj* in the type hierarchy and represents, essentially, a type error. The special type  $\perp$  is at the bottom of the type hierarchy and is used in initialization of an analysis.

$$\overline{A <: \top} \quad , \quad \overline{\perp <: A}$$

There are no generics in Quack, neither through a template facility for user-defined generic classes, nor through built-in generic facilities like Java arrays. Quack also has no function types (i.e., functions are not first-class values), and no nested functions or inner classes. All the hard stuff has been omitted.

Although the concrete syntax of Quack includes expressions, arithmetic expressions are syntactic sugar for method calls (e.g.,  $a + b$  is syntactic sugar for  $a.\text{plus}(b)$ ). This reduces the number of cases that must be accounted for in the type-check rules.

Unlike arithmetic expressions involving '+', '-', '\*', and '/', the boolean operations *and*, *or*, and *not* are not syntactic sugar for method calls. Interpreting these as method calls would not permit short-circuit evaluation. This is a consequence of call-by-value semantics: Actual arguments are evaluated before being passed to a method. A method *and* that took two arguments would receive the already evaluated results of its left and right-hand operands, and would not be able to prevent evaluating the right-hand operand if the left-hand operand was False. For example,  $(\text{not } (x == 0))$  and  $(1/x < 5)$  should never cause a divide-by-zero error, but it could if *and* was a method.<sup>2</sup>

## 2.1 Explicit type declarations

Explicit type declarations appear in the formal argument lists of class constructors (given in the class header) and formal argument lists of methods:

```
class Pt(x: Int, y: Int) {
  ...
  def PLUS(other: Pt): Pt {
    return Pt(this.x + other.x, this.y + other.y);
  }
}
```

An explicit type declaration may optionally appear in an assignment:

```
a_square: Rect = Square( Pt(3,3), 5 );
```

Where explicit type declarations appear, they determine the types of the variables they name. In the type inference procedure described below, we will treat type declarations as if they appeared separately from the assignment statement.

<sup>2</sup>If you are using Haskell, you may permit yourself a moment of gloating because it is perfectly possible in a *lazy* language to define functions that are *non-strict*, that is, functions that evaluate arguments only if and when their value is actually used. However, adding laziness to Quack would cost us a good deal in other complexity.

## 2.2 Instance variables

Instance variables (also called ‘data members’ or ‘fields’ of an object) are introduced by assignment in the constructor part of a class declaration. There is no separate variable declaration, nor explicit indicator of the type, although an explicit type may be included in an assignment to an instance variable just as in other assignments. The static type of the instance variable is inferred from the expression whose value is assigned to it. Consider:

```
class Square(ll: Pt, side: Int) extends Rect {
  this.ll = ll;
  this.ur = Pt(this.ll._x() + side, this.ll._y() + side);
}
```

In class Square, instance variable `this.ll` has (inferred) static type `Pt` because formal argument `ll` has (declared) static type `Pt`. Instance variable `ur` has static type `Pt` because the constructor for class `Pt` returns a `Pt` object.

If an explicit type declaration is included in an assignment, it determines the static type. If the declared type is the same as the inferred type, it has no effect. If the inferred type is not a subtype of the declared type, a type error is reported. For example:

```
class Square(ll: Pt, side: Int) extends Rect {
  this.ll: Obj = ll;    // OK, Pt <: Obj
  this.ur = Pt(this.ll._x() + side, this.ll._y() + side);
  // Error: No method _x() in class Obj
}
```

Control flow is possible in a constructor. The set of instance variables defined in the constructor must be independent of the control flow. That is, the following is *not* permitted:

```
class Schroedinger(box: Int) {
  if box > 0 {
    this.living = 1;
  } else {
    this.dead = 0;
  }
}
```

However, the following is permitted:

```
class Schroedinger(box: Int) {
  if box > 0 {
    this.living = true;
  } else {
    this.living = false;
  }
}
```

In the second version, while the value of `this.living` depends on control flow, it has been initialized with and has static (and dynamic) type `Boolean` in either case. But this then raises a question as to whether the following is permitted, and if so what types should be inferred:

```
class Hand() {
    /* Nothing to see here */
}

class LeftHand(x: Int) extends Hand {
    this.x = x;
    def foo(): Int { return 42; }
}

class RightHand(x: Int) extends Hand {
    this.x = x;
    def foo(): Int { return 42; }
}

class Bot(x: Int) {
    if x > 0 {
        this.hand = LeftHand(3);
    } else {
        this.hand = RightHand(7);
    }
}
```

This is permitted, and gives `this.hand` the static type `Hand` (the nearest common ancestor of `LeftHand` and `RightHand` in the class hierarchy). Since the static type `Hand` is a class that does not have the `foo` method, changing the `Bot` method as follows will cause it to fail with a type error:

```
class Bot(x: Int) {
    if x > 0 {
        this.hand = LeftHand(3);
    } else {
        this.hand = RightHand(7);
    }
    this.answer = this.hand.foo(); // Nope, no foo in Hand
}
```

## 2.3 Inherited Instance Variables

Any instance variable in a superclass must be introduced in the constructor of all subclasses with compatible types.

```

/**
 * Inheriting instance variables --- must be
 * consistently initialized.
 */
class Robot() {
    this.name = "Bozotron";
    this.age = 42;
    this.strength = 100;

    def destroy_city(fierceness: Int ) {
        /* TBD: how do we actually destroy things? */
    }
}

class SmartRobot(iq: Int) extends Robot {
    this.name = "Einsteintron"; // OK, compatible
    this.age = "forty-eight";   // ERROR, not an Int
    this.iq = 2 * iq;           // Sure, why not
    // ERROR, failed to define strength
}

```

This rule is necessary because the inherited method `destroy_city` might use its strength, and method `Int.PLUS` might be applied to its age while our robot is imprisoned.

## 2.4 Method Signatures

Formal arguments of other methods are treated like formal methods of class constructors: They have explicitly declared static types that cannot be changed. The static types of instance variables (which may be introduced only in the constructor) are fixed in methods. That is, methods do not alter the static types of instance variables, although they may assign to an instance variable any compatible value including values belonging to any subtype of the instance variable's static type.

If a method in a subclass has the same name as a method in a superclass, we say it *overrides* the superclass method. Dynamic dispatch may result in calling the subclass method where the superclass method is expected, so we must ensure the overriding method is compatible with the superclass method. Thus:

```

class Pt(x: Int, y: Int) {
    this.x = x;
    this.y = y;

    def STR() : String {
        return "(" + this.x.STR() + ","
            + this.y.STR() + ")";
    }
}

```

```

def PLUS(other: Pt) : Pt {
    return Pt(this.x + other.x, this.y + other.y);
}

def _x() : Int { return this.x; }
def _y() : Int { return this.y; }
}

class Rect(ll: Pt, ur: Pt) extends Obj {
    this.ll = ll;
    this.ur = ur;

    def translate(delta: Pt) : Pt { return Rect(ll+Pt, ur+Pt); }

    def STR() {
        lr = Pt( this.ur._y(), this.ll._x() ); // lower right
        ul = Pt( this.ll._x(), this.ur._y() ); // upper left
        return "(" + this.ll.STR() + ", "
            +      ul.STR() + ", "
            +      this.ur.STR() + ", "
            +      lr.STR() + ")";
    }
}

class Square(ll: Pt, side: Int) extends Rect {
    this.ll = ll;
    this.ur = Pt(this.ll._x() + side, this.ll._y() + side);
}

a_square = Square( Pt(3,3), 5 );
a_square = a_square.translate( Pt(2,2) );
a_square.PRINT();

```

In the example above, STR is a method inherited from Obj, like the toString method in Java.

PLUS is not inherited from a built-in class, but any class may define a PLUS method to take advantage of syntax for the arithmetic operator +.  $x + y$  is syntactic sugar for  $x.PLUS(y)$ . If a compatible PLUS method does not exist for the static type of  $x$ , the Quack compiler will report an error, just as if the programmer had written  $x.PLUS(y)$  explicitly.

The full rule for compatibility of an overriding class is that

- The number of formal arguments of the overriding method must be the same as the number of arguments in the method it overrides.

- For each formal argument  $a$  of the overriding method, and the corresponding argument  $a_{\text{inherited}}$  of the method it overrides,  $a_{\text{inherited}} \leq a$  (the overriding method must accept at least everything the inherited method would have accepted, if it had not been overridden).
- The type  $r$  returned by the overriding method must be a subtype of the type  $r_{\text{inherited}}$  of the superclass method.

## 2.5 Local Variables

This leaves the static types of local variables to consider. The rules are essentially like those for instance variables, except that it is not necessary that every execution path instantiate the same set of instance variables, only that an instance variable is initialized before it is used on every execution path.

### 2.5.1 Initialization before use

A variable must be initialized before it is used on every *syntactic* execution path. The following is not permitted:

```
def bad_init() {
    rep = 0;
    n_reps = 3;
    while rep < n_reps {
        if rep > 0 {
            x = y;    // Error here
        }
        rep = rep + 1;
        y = 42;
    }
}
```

In the code above, it is not possible for the assignment to  $x$  to occur before the assignment to  $y$  on any feasible execution path, because the condition  $\text{rep} > 0$  will not be true on the first iteration of the loop. Nonetheless this program will be rejected by the Quack compiler because of a potential use of an uninitialized value (reference before definition) on the syntactic path that enters the loop and takes the true branch of the conditional. Essentially, for consideration of initialization, we treat every boolean condition as if it could be either true or false.

This restriction, together with the absence of a way to create a null pointer, allows us to avoid dealing with null pointers in execution.

The remaining problem is to infer the types of local variables, which we consider in the next section.

## 3 Inferring Local Variable Types

### 3.1 A Constraint from Concrete Semantics

The basic property that static type inference must ensure is that the types of all the values a variable may hold are compatible with the static type of that variable. That is because we will use the static type to determine what methods may be called on objects held by that variable. Really bad things would happen if we called `x.foo(1, "yup")` on an object that does not have a `foo` method that takes an integer and string as arguments. In a dynamically typed language, we check consistency every time a method is called, and fail gracefully (perhaps by crashing the program) if there is no such method. In a statically typed language, we do not make that check at run-time.<sup>3</sup>

A strongly typed, statically typed language like Quack *might* just crash if a variable held a value that did not conform to its static type, but it could be much worse. It will execute the method even *assuming* that all the values are of the correct type, which could result in arbitrary behavior, like the possible behavior of a C or C++ program in which an integer has been cast to a function pointer and then called. This is known as “catch fire” semantics (it is not a violation of the language semantics if your computer catches fire), and is the stuff of security nightmares. We’d like to prevent that.

There is a simple way to ensure that the static type of a variable is compatible with all the object values it might hold: We could use ‘Obj’ as the static type for everything. Great, problem solved. Unfortunately, we can’t write any useful programs, because ‘Obj’ doesn’t have the methods we’ll want to call on those objects. We can say that ‘Obj’ is a safe approximation of the static type we want (it will not cause our computer to catch fire), but it is *too coarse* an approximation. Our goal should be to find as precise an approximation as possible, while still being safe.

In terms of the class hierarchy and the corresponding hierarchy of types, we want to find an approximate type that is a supertype of all the types of all the values that a variable may hold at run-time, but no higher in the hierarchy than necessary.

### 3.2 A Collecting Semantics

To find the best approximation for the set of actual object types a variable may hold at run time, we can imagine an alternative way of executing the program in which, instead of holding actual values, variables hold just indicators of the types of objects. If we did this, we would not be able to determine whether the true branch or the false branch should be taken by an if statement, nor how many times a loop executes, so we would need to consider all syntactic paths (as we did when considering initialization) rather than just the execution paths that can actually be taken.

---

<sup>3</sup>Mostly we do not make that check at run-time. Some statically typed languages make a dynamic type check in a few places where their static type system cannot ensure consistency. For example, Java makes a dynamic type check when you cast an Object value to something more specific, and when you access an element of an array. The restrictions in Quack, including lack of built-in or user-defined generic collections, make it unnecessary for us to make run-time type checks.



If the static types given to variables were safe when considering all these paths, then it would also be safe when considering just the paths that can actually be executed.

The number of execution paths through a program is potentially infinite, which makes it difficult to enumerate them all. However, this abstract version of the execution has an important advantage over the normal program semantics: There are only a finite set of types in a Quack program, so if we keep enumerating paths through a method, eventually we will encounter only states that have been explored before.<sup>4</sup>

This would be enough to define a sound but reasonably precise approximation to use as the static type of a local variable: We could enumerate all the possible types of all the objects assigned to a variable on every syntactically legal path, and find their least common ancestor in the class hierarchy. It's a start.

### 3.2.1 A Note on Modular Checking

Note that Quack requires explicit type declarations on method arguments (including constructor arguments) and on the return type of a method. This is so that, whether we are using the collecting semantics or a more practical abstract semantics introduced below, we never have to actually execute a method call. We can just use the method signature of the current method to determine the types of arguments passed to the current method, and we can use the signature of called methods to determine whether the method call is legal and what type of value it will return. Instance variables are fixed by the constructor for the same reason: We must analyze the constructor before we analyze other methods in a class, but after that we can consider each method independently, without worrying about how types in one affect types in another. The requirements for explicit declarations in Quack are there to permit you to perform type inference and type checking in a modular fashion, one method at a time.

### 3.2.2 Execution steps and states

Let us consider what an execution step would be with the collecting semantics. We can imagine that each complex expression has been transformed into a series of assignments to temporary variables, so that for example

```
x = y + z.foo(w,q);
```

has been expanded to

```
t1 = z.foo(w,q);  
x = y.PLUS(t1);
```

Thus broken down into small pieces, the execution steps we need to consider are method calls (including determination of the type of the receiver object, such

---

<sup>4</sup>While an individual method is *finite state* in this semantics, a complete Quack program may not be, because of potential recursive method calls. Fortunately Quack also has a very restricted scope mechanism which, together with the requirement that method declarations declare the types of their arguments and return value, permits us to consider each method body in isolation. This is also why we insist that object instance variables be given a static type in the constructor that is not altered by methods.

as  $y$  in the example above), assignment, boolean expressions, and the control flow constructs *if* and *while*.

We will model the environment of a computational step as a map from variables to sets of value types. Initially, a variable with a declared type  $T$  will be mapped to the singleton set of types  $\{T\}$ . All other variables will be initially mapped to the empty set of types.

**Expressions: Method calls** Consider evaluation of the expression  $z.m(a_1, a_2, \dots)$  in an environment that has a non-empty mapping for  $z$  and for each  $a_i$  in the list of actual arguments passed to method  $m$ . That is, the environment in which the expression is evaluated looks like this:

$$(z \mapsto \{z_0, z_1, \dots\}, a_1 \mapsto \{a_1^1, a_1^2, \dots\}, a_2 \mapsto \{a_2^1, a_2^2, \dots\})$$

To determine the set of types that evaluation of this method call might return, we simply consider every possible combination of types for the variables. For each type  $z_i$  that  $z$  might have, we consider method  $m_i$  of class  $z_i$ . The declared return type of  $m_i$  is one of the possible types of the expression.

**Type errors.** What if some  $m_i$  has the wrong number of formal arguments, or if one of the potential types of an actual argument  $a_j^i$  is not compatible with the corresponding formal argument of  $m_i$ ? In that case we have detected a potential type error. We will not be able to complete the type inference, but must report the potential error instead. Formally we could define a special  $\top$  and give the expression that type, but if we continue with type inference with this “error” type we will soon generate a large number of annoying error messages; it is likely best to give up while we’re ahead.

**Uninitialized variables.** What if one or more of the the variables  $z$  or  $a_i$  has an empty set of possible types? We can ensure this never happens. Recall that Quack has a separate rule that requires every variable to be initialized before it is used on every syntactic program path. As long as we consider expressions in some order that corresponds to a potential execution path, we are assured that each variable we use in type inference is mapped to a non-empty set of possible types. The lexical order in which statements appear in the program may not correspond to any possible execution order (for example, statements in the ‘else’ branch of an ‘if’ statement cannot actually be executed immediately after statements in the ‘then’ branch), but it is also certain to be an acceptable order for purposes of type inference.

**Boolean expressions.** Boolean operations *and*, *or*, and *not* are not syntactic sugar for method calls, but for the purpose of type checking and type inference we can treat them as if they were. Short circuit evaluation is irrelevant to type checking, so we can pretend that these operations are shorthand for methods of class Boolean<sup>5</sup>

---

<sup>5</sup>Treating these boolean connectives as method calls is useful for explanation, but in your implementation it may be simpler to just implement the type checking for these connectives directly.

We'll use names that are illegal in Quack so that it is impossible for them to conflict with any method names that a Quack programmer could create:

```
class Boolean() {  
    ...  
    def $AND(Boolean other): Boolean { ... };  
    def $OR(Boolean other): Boolean { ... };  
    def $NOT(): Boolean { ... };  
}
```

**Assignment statements.** In the collecting semantics, an assignment statement *adds to* the set of types that a variable might hold. Suppose we have the statement

```
x = z.foo(w,q);
```

Before the statement, the environment is  $(\dots, x \mapsto S_1, \dots)$ , and we find that in that environment the expression `z.foo(w,q)` can have some set of types  $S_2$ . Since this is a *collecting* semantics, we do not replace the old value  $S_1$  but rather augment it: The new environment is  $(\dots, x \mapsto S_1 \cup S_2, \dots)$ .

**Control flow and iteration to a fixed point.** How should we interpret the control flow statements *if/elif/else* and *while* when we are keeping track only of the types of variables, and not their values? How can we know which branch of an *if* statement to take, or how many times to execute a loop? The answer is that we do not have to decide. In the collecting semantics, we consider all possible paths through the control flow (syntactic paths again!).

We do need to consider evaluation of the boolean expressions in *if* and *while* statements, but only to catch type errors. The answer to “which way should we go” will always be “both ways.”

How can this possibly work? There are an infinite number of paths of infinite length in any method containing even a single *while* loop; we cannot trace them all to the end. However, there is only a finite number of types (classes) in a Quack program. When we modify the program state in this collecting semantics, it is always by adding types to the potential types of a variable. Thus, at some point on every path, we must encounter exactly the same assignment of sets of types.

In fact, as long as we make the first pass through in *some* legal program order to avoid dealing with uninitialized variables, it doesn't even matter what order we evaluate statements. We can simply run through all statements in their lexical order again and again until we make a complete pass with no changes to the execution state. When this happens, we say that we have reached a *fixed point* in the evaluation, and that fixed point a safe approximation of the set of actual object types that can be assigned to each local variable.

**Static types.** The static type of a variable can be chosen as the least common ancestor of all the set of types of the objects it can hold. It should be clear that this choice is a safe estimate.

What may be less clear is that it is sufficiently precise. For example, suppose at the conclusion of simulating the collecting semantics, we had a variable  $x$  mapped to the set of types  $\{\text{Int}, \text{String}\}$ . The common ancestor of *Int* and *String* in the type hierarchy is *Obj*. *Int* has a method *PLUS* that adds two integers, and *String* has a method *PLUS* that concatenates two strings. However *Obj* has no method *PLUS*. It seems, therefore, that we might choose a static type for which the *PLUS* method (and its sugared form  $+$ ) would be disallowed despite all uses of *PLUS* being applied correctly. It is easy to construct a program that would be accepted by a dynamic type system but is rejected by our static type system for exactly this reason:

```
def not_duck_typing(x: Int): Int {
  if x < 7 {
    a = 42;
    b = 13;
  } else {
    a = "forty-two";
    b = "thirteen";
  }
  if a < b {
    return 1;
  } else {
    return 2;
  }
}
```

If the actual argument passed as  $x$  is less than seven,  $a$  and  $b$  will both be integers, and we might expect the method to return 2. If the actual argument is seven or greater, we might expect it to return 1, because “forty-two” is before “thirteen” in collating order (the meaning of comparison for strings). That is what Python, with dynamic types and a “duck typing” rule, would do. But that is not what Quack or other statically typed languages do. By “static” we mean that the types are determined at compile time, so we need to know whether  $a < b$  is a numeric comparison or a string comparison. The common ancestor of *Int* and *String* in the type hierarchy is *Obj*, and *Obj* has no method corresponding to the  $<$  comparison operation.

### 3.3 An Abstract Semantics

The collecting semantics is a little unwieldy, and it is also clumsy to find a fixed point solution on sets of types, and then take the least common ancestor of each set, and finally discover a type error. We can address both problems by taking the least common ancestor on-the-fly, using it as a representation for the set of types a variable could take.

There is no natural representation for the empty set of possible types, but we can augment the type hierarchy with a special element  $\perp$  (pronounced “bottom”) to represent “no value yet”. We can also add an element  $\top$  (pronounced “top”) representing a type error. If we place these at the bottom and top of the hierarchy as shown in Figure 2, the types form a lattice ordered by the “subtype of” property.

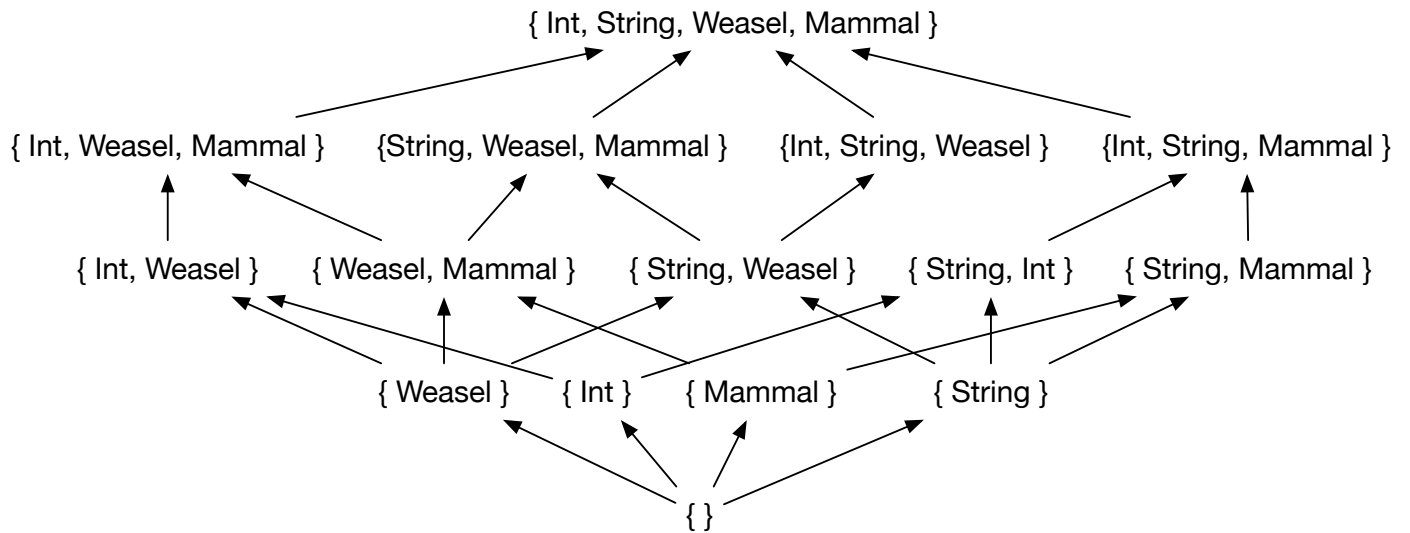


Figure 1: Subsets of types form a semi-lattice ordered by set inclusion. In the example, Weasel is a subclass of Mammal.

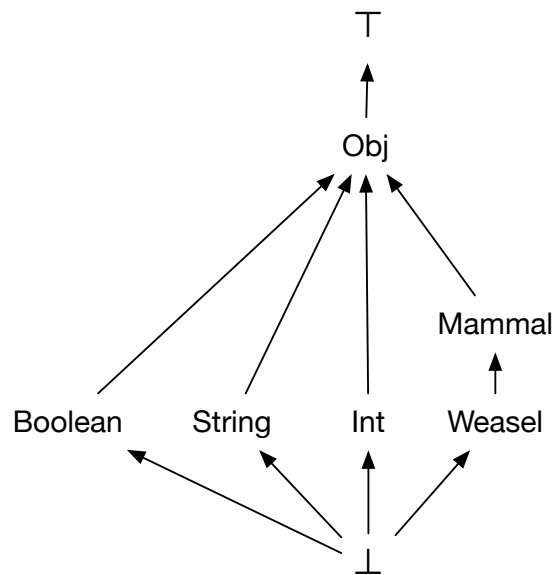


Figure 2: Types ordered by the “subclass of” relation, augmented with an element  $\perp$  for “no information yet” and  $\top$  for “type error”, form a lattice. Compare to Figure 1. Weasel is a subclass of Mammal.

Like the lattice of sets ordered by “subset of”, the lattice of types ordered by “subtype of” has a property that is useful in reasoning about correctness and termination of our type inference. Suppose we have an expression  $z.o(a_1, a_2, \dots)$  that gives us a type  $T$  in environment  $E$ . Now suppose we evaluate the same expression in another environment  $E'$  which is identical except that some of the arguments  $a'_i$  are supertypes of corresponding  $a_i$  in  $E$ . It is not too hard to see that  $T' \geq T$ , i.e., when the arguments to a method (including the receiving object) go “up” in the lattice, the resulting type either stays the same or goes up; it cannot go down or “across”. (We added the  $\top$  element at the very top so that type errors would not be an exception to this rule.) We say that the type inference rule is *monotone*. Being monotone, it ensures that the computation of inferred static types converges (reaches a fixed point) on the types we would have obtained using the sets of types directly. In fact it does so very rapidly, because the class hierarchy in an object-oriented program is almost always quite shallow.

**We didn’t want duck typing, anyway.** It might seem at first that rejecting method `not_duck_typing` is a flaw in our type inference, or at least a disadvantage. However, using the least common ancestor rather than the set of possible types has another advantage in addition to making type inference and type checking easier: It allows us to use a common implementation strategy for dynamic method dispatch. All the subclasses of a class with a method *foo* are not only guaranteed to have a compatible method *foo* (either inherited or overridden), but in addition we will implement those classes with a pointer to the code for *foo* in the same position in the class description. The table of pointers to methods in the class is sometimes called the *vtable*, because in C++ the methods in that table are called “virtual functions” or “virtual methods.”

### 3.4 Putting it Together: Implementing Type Checking

The abstract semantics above is the one you should implement. You will need an actual representation of the subtype hierarchy for all the subtypes of *Obj*. The  $\top$  and  $\perp$  elements can be implicit (which simplifies matters since the rest of the hierarchy is a simple tree that can be represented by a table).

You will want to create a table that maps from variables visible in a method to the types of those variables. Initially the arguments to the method have their declared types, and all other variables have type  $\perp$ .

In pseudocode, the logic of the main type inference loop for a single method in Quack is roughly this, if

```

changed = true;
while changed {
    changed = false;
    for each statement, in lexical order {
        type check the statement
    }
}

```

Of course I've hidden some complexity in the “type check the statement” step. Really it's not so complex, but each kind of Quack statement and expression can require its own code. We handle the type of an expression exactly as we would do for the collecting semantics, except that now we don't have to deal with sets of possible types, but just with the current estimate of a type.

For assignment statements, we don't just replace current type estimate for a variable. Rather, if we have previously estimated the type of a variable to be  $T$ , and we now we are assigning type  $Q$  to that variable, the new type is the *least common supertype* of  $Q$  and  $T$  (i.e., their nearest ancestor in the tree). Often, but not always, that will be either  $Q$  or  $T$ . Note that our estimate may “move up” the tree but never down or across; and since the tree is finite, this is our guarantee that type inference will terminate. If this results in a new type estimate for the variable, we set the ‘changed’ flag that controls the outer loop in the pseudocode above.

What if the assignment statement has a declared type, like

```
x: Int = a + b;
```

After inferring a type for  $a + b$ , we check that the inferred type is a subtype (possibly equal). If not, we report a type error. Then we set the estimated type to the declared type (even if it is a proper supertype of the inferred type).

**Other statements.** Besides assignments, Quack has while loops, if statements, return statements, raw expressions, and typecase statements.

- For while loops, we check that the condition evaluates to Boolean, and then we type check each statement in the body of the loop in the normal way.
- For if statements, we check that the condition evaluates to Boolean, and then we type check each statement in the true branch and the false branch.
- For return statements, we check that the expression evaluates to a some subtype of the method return type. (In some cases this will be “Nothing”.)
- For an expression that appears by itself, we calculate the type of the expression (because we might find a type error in the evaluation) and then discard it.
- Typecase is the most complex of the cases, and is considered below.

**Typecase** A typecase statement is used to choose different code depending on the dynamic type of an object, which may be a proper subtype of the static type of the variable that holds it. Consider:

```
def EQUAL(other: Obj) : Boolean {  
  typecase other {  
    pt: Pt { return this.x == pt.x and this.y == pt.y; }  
  }  
  return false;  
}
```

We evaluate the typecase expression (*other* in this example) and discard the result; we only care that it can be evaluated without a type error. Then, for each case (of which this example has only one), we introduce a variable (like *pt*) with the declared type (*Pt* in this case), and type check the statements within that branch of the typecase in the augmented environment. (Note that the scope of the introduced variable is only for that branch of the typecase.)